

Offensive AI

A fun and not (so) explored field mixing **offensive cybersecurity** and
machine learning.



About Me!

Dorian Bachelot

Research Engineer

*@LHS/CentraleSupélec/INRIA
#Malware Developer/Detection*

<https://dorianb.net>



Introduction

AI, ML, DL, OAI...

Field Overview

Field Overview

Definition

Offensive artificial intelligence (OAI) can be defined as the field in which artificial intelligence algorithms or models are used to **automate one or multiple parts of an offensive security task**.

Field Overview

Definition

Offensive artificial intelligence (OAI) can be defined as the field in which artificial intelligence algorithms or models are used to **automate one or multiple parts of an offensive security task**.

This cover:

- Vulnerability scanning
 - Exploit automation
 - Attack path identification
 - ...

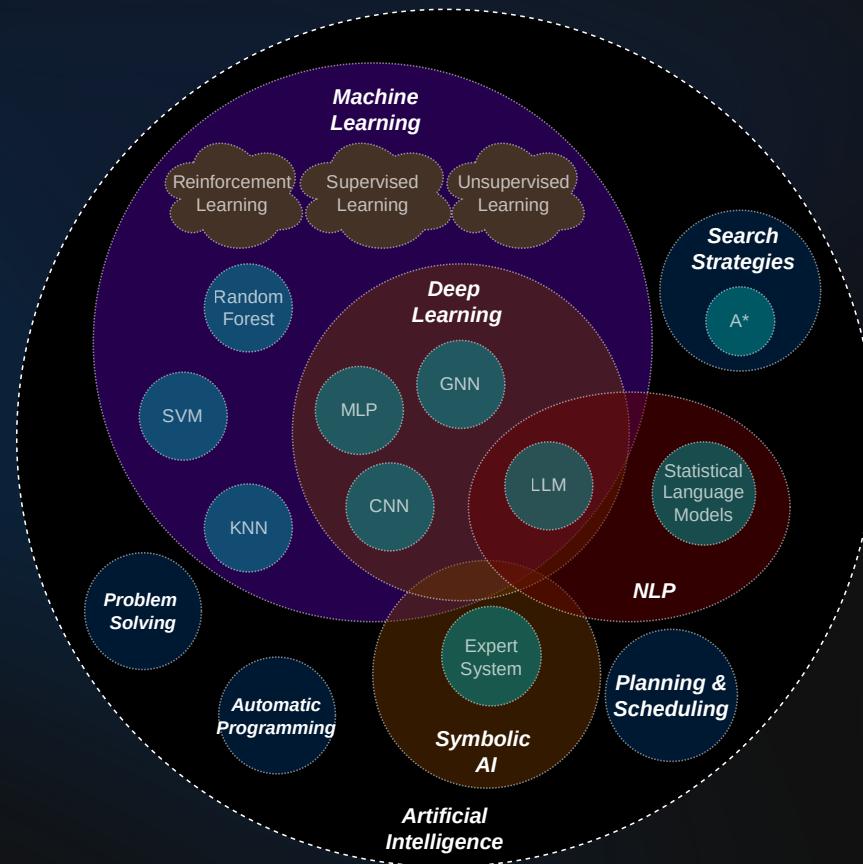
Reminder on AI

Reminder on AI

We can classify as AI any system that solves a problem or achieve a goal and have some degree of adaptation and autonomy.

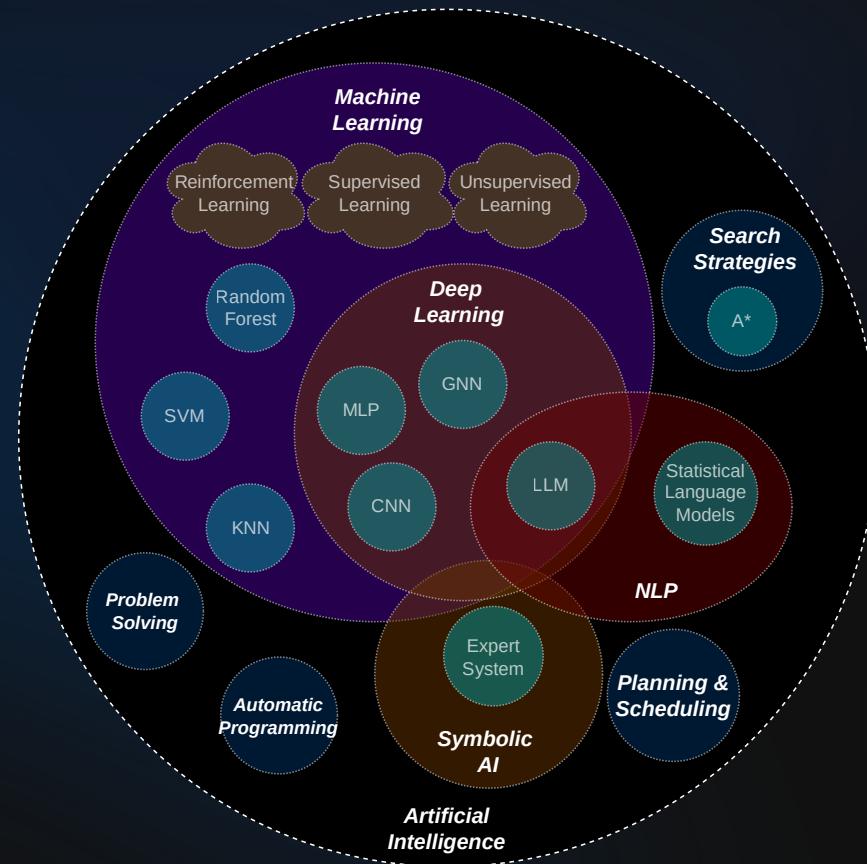
Reminder on AI

We can classify as AI any system that solves a problem or achieve a goal and have some degree of adaptation and autonomy.



Reminder on AI

We can classify as AI any system that solves a problem or achieve a goal and have some degree of adaptation and autonomy.



Machine Learning \neq AI \neq Deep Learning \neq LLMs

OAI

Global overview

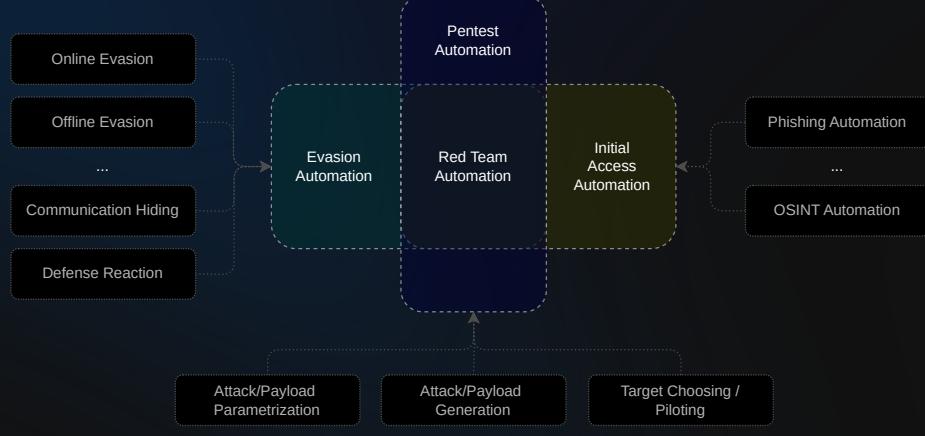
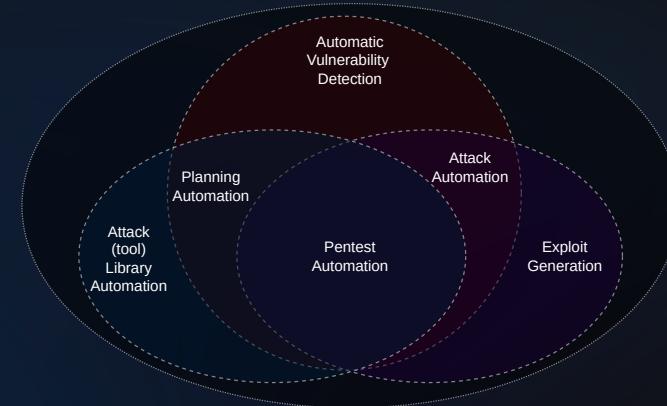
Overview

Overview

OAI can also be cut in **multiple subfields**, which correlate closely to some **use cases** like **pentest** or **red team automation**.

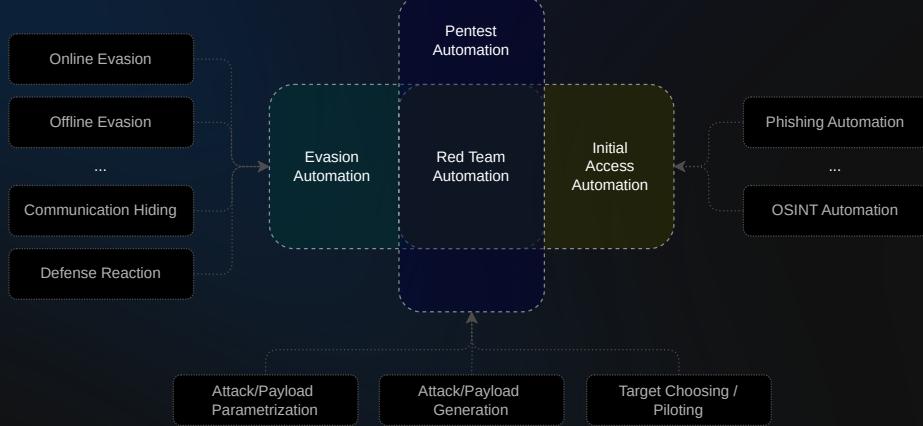
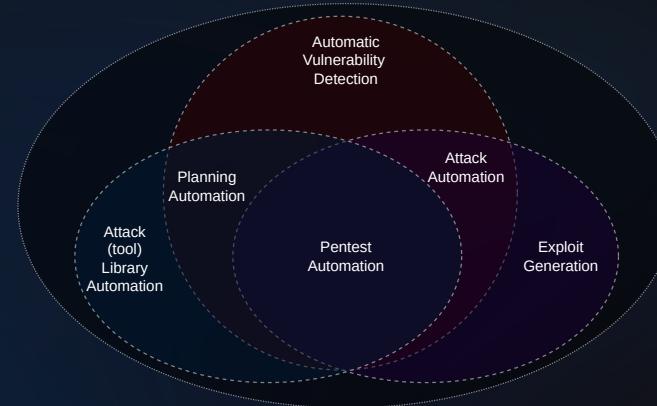
Overview

OAI can also be cut in **multiple subfields**, which correlate closely to some **use cases** like **pentest** or **red team automation**.



Overview

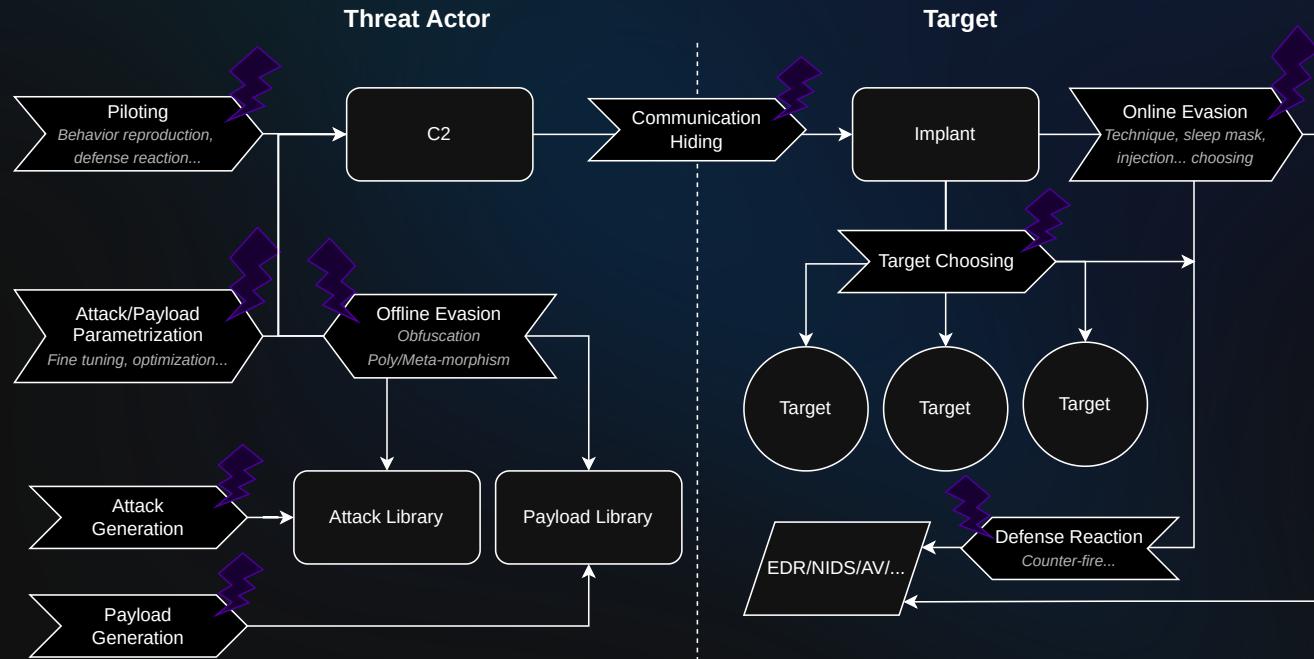
OAI can also be cut in **multiple subfields**, which correlate closely to some **use cases** like **pentest** or **red team automation**.



Automation ≠ Better than Humans

The different subfields

The different subfields



OAI for C2 Piloting

Target choosing, attack choosing...

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

○●oooooooooooooo

OAI for Evasion

oooooooooooooooo

Conclusion

oooo

C2 Piloting

C2 Piloting

What is a C2?

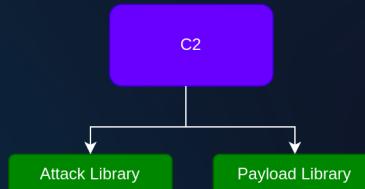
- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and send attacks (modules).
- Focus on evasion, hiding communication and using advanced implants.
- Used to pivot and deeply infect a network.
- Controlled by an operator (or AI!).

C2 Piloting

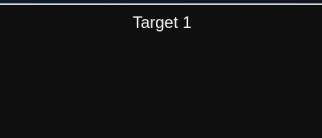
What is a C2?

- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and send attacks (modules).
- Focus on evasion, hiding communication and using advanced implants.
- Used to pivot and deeply infect a network.
- Controlled by an operator (or AI!).

Threat Actor



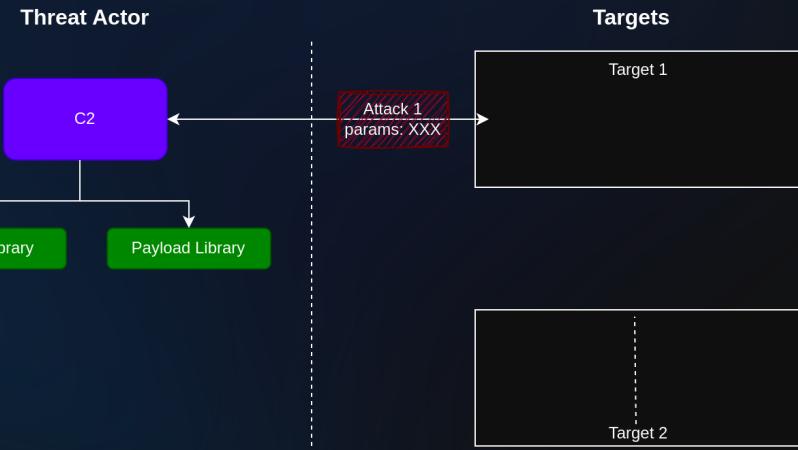
Targets



C2 Piloting

What is a C2?

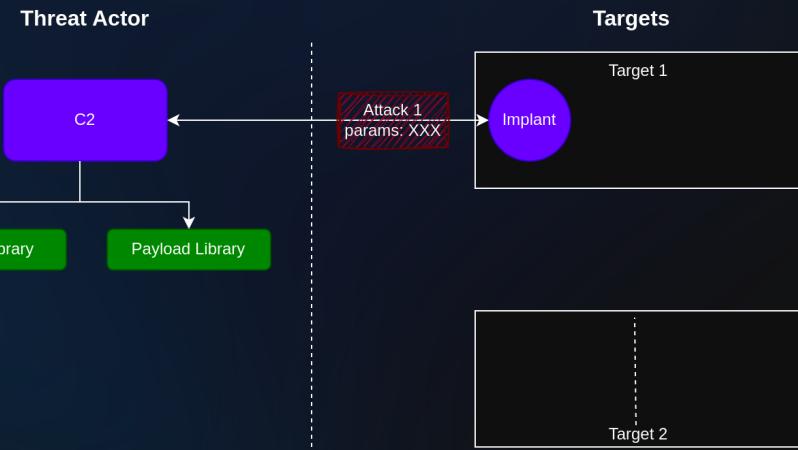
- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and send attacks (modules).
- Focus on evasion, hiding communication and using advanced implants.
- Used to pivot and deeply infect a network.
- Controlled by an operator (or AI!).



C2 Piloting

What is a C2?

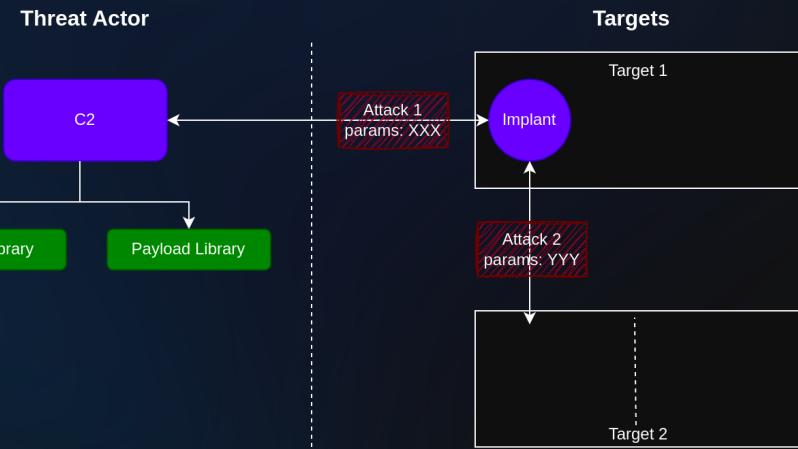
- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and send attacks (modules).
- Focus on evasion, hiding communication and using advanced implants.
- Used to pivot and deeply infect a network.
- Controlled by an operator (or AI!).



C2 Piloting

What is a C2?

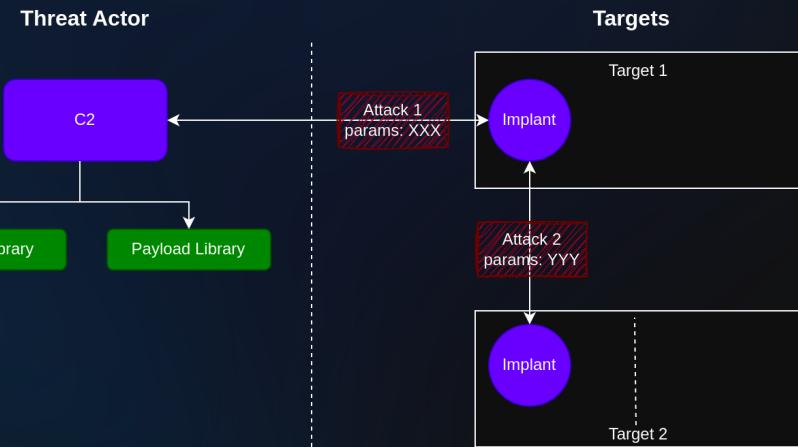
- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and **send attacks** (modules).
- Focus on **evasion**, hiding communication and using advanced implants.
- Used to **pivot** and **deeply infect** a network.
- Controlled by an operator (or AI!).



C2 Piloting

What is a C2?

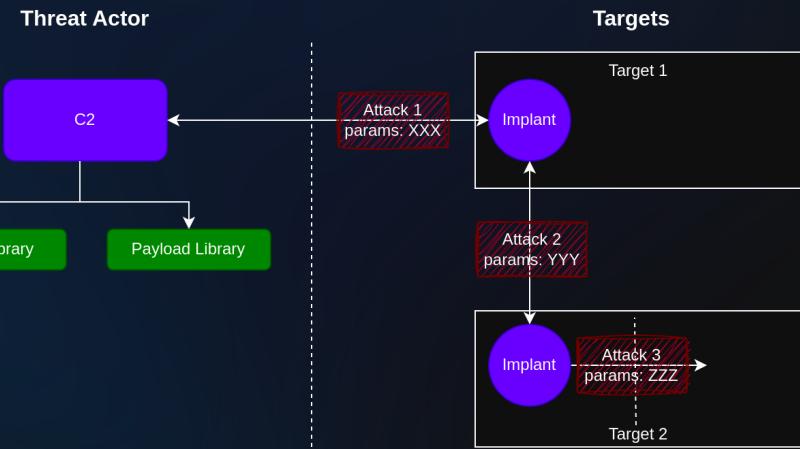
- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and **send attacks** (modules).
- Focus on **evasion**, hiding communication and using advanced implants.
- Used to **pivot** and **deeply infect** a network.
- Controlled by an operator (or AI!).



C2 Piloting

What is a C2?

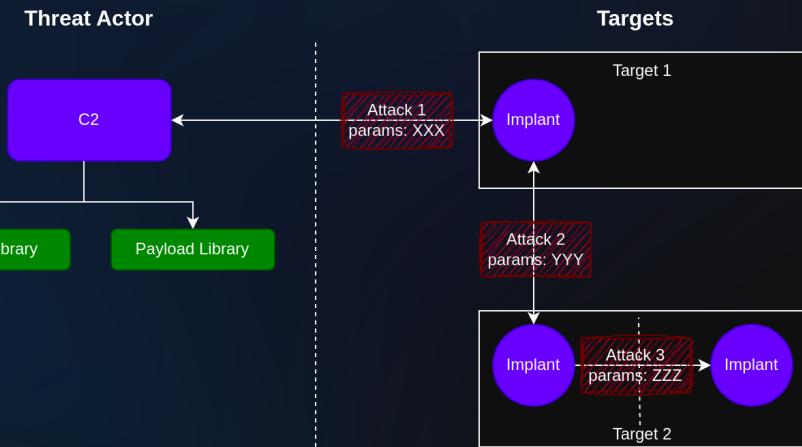
- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and **send attacks** (modules).
- Focus on **evasion**, hiding communication and using advanced implants.
- Used to **pivot** and **deeply infect** a network.
- Controlled by an operator (or AI!).



C2 Piloting

What is a C2?

- Pilot implants, that have infected some hosts.
- Implants can communicate directly with the C2, or can do it through another implant.
- Can issue commands, and send attacks (modules).
- Focus on evasion, hiding communication and using advanced implants.
- Used to pivot and deeply infect a network.
- Controlled by an operator (or AI!).



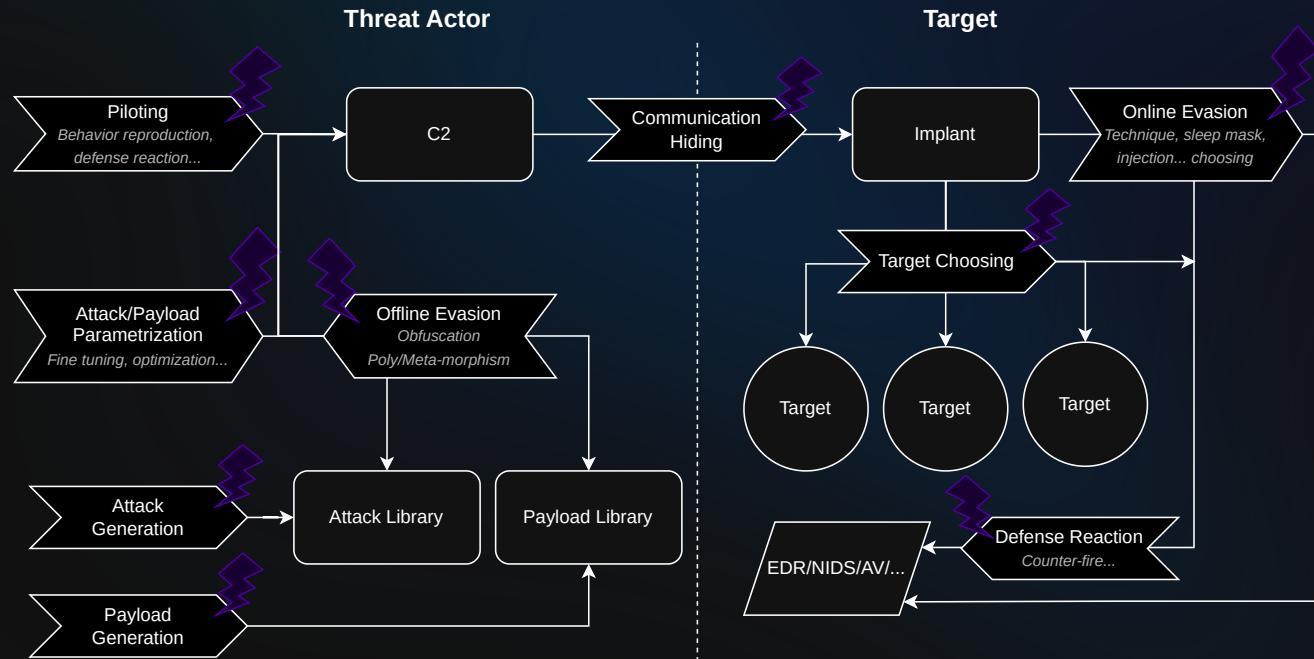
C2 Piloting / The Idea

C2 Piloting / The Idea

Consists of selecting a **target**, an **attack** and **parameters** using a model/algorithm.

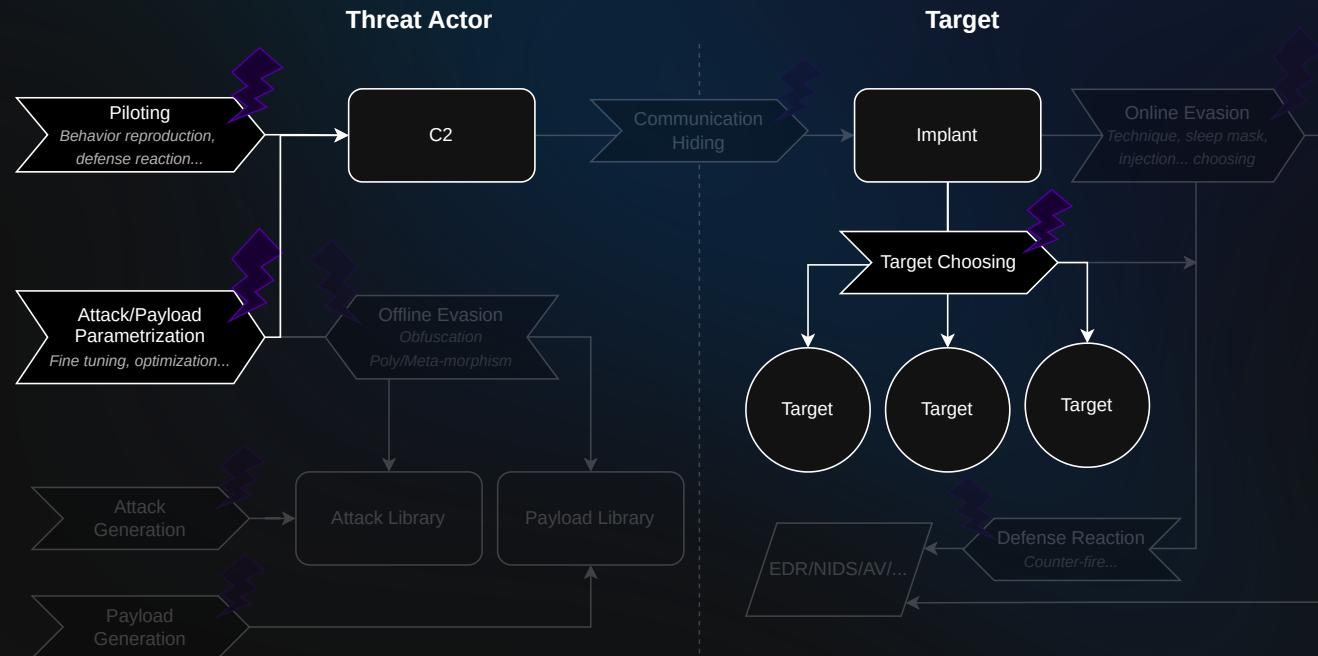
C2 Piloting / The Idea

Consists of selecting a **target**, an **attack** and **parameters** using a model/algorithm.



C2 Piloting / The Idea

Consists of selecting a **target**, an **attack** and **parameters** using a model/algorithm.



Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooo●oooooooooooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

C2 Piloting / What's Needed?

C2 Piloting / What's Needed?

Recipe:

- Take a **C2 with an API** (giving automation power on ALL operations).
- Prepare an **attack library** (modules), with exposed parameters.
- Add some model/algorithm.
- Shake. Done!

C2 Piloting / What's Needed?

Recipe:

- Take a C2 with an API (giving automation power on ALL operations).
- Prepare an attack library (modules), with exposed parameters.
- Add some model/algorithm.
- Shake. Done!



(@chrisrohlf)

C2 Piloting / What's Needed?

Recipe:

- Take a C2 with an API (giving automation power on ALL operations).
- Prepare an **attack library** (modules), with exposed parameters.
- Add some model/algorithm.
- Shake. Done!



(@chrisrohlf)

Not that simple in reality... A lot of AI-related issues.

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooo●oooooooooooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

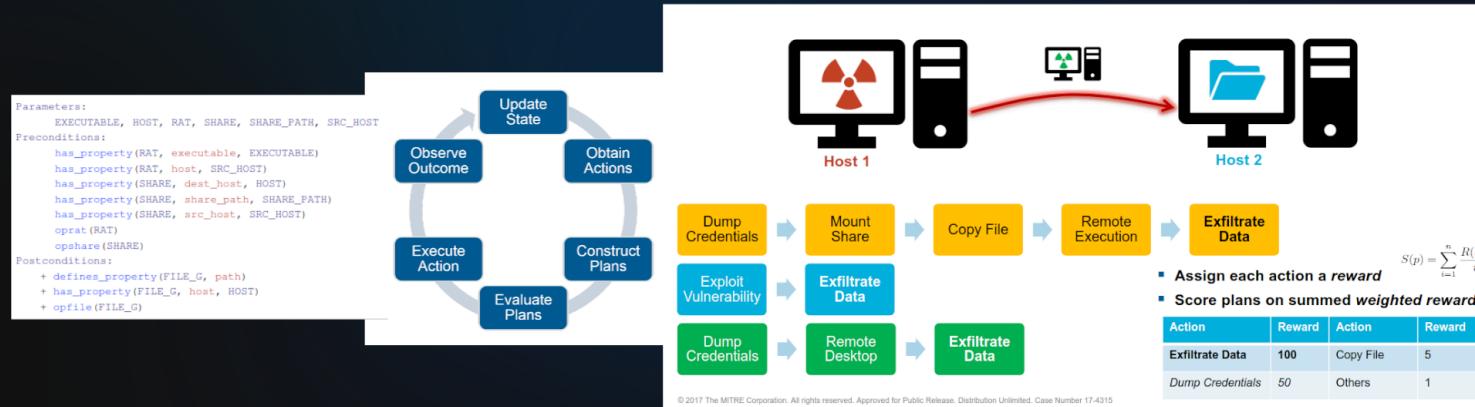
C2 Piloting / Heuristics

C2 Piloting / Heuristics

Caldera⁽¹⁾, an automated framework for adversary emulation (by Mitre), rely on a **heuristic** to select attacks & targets.

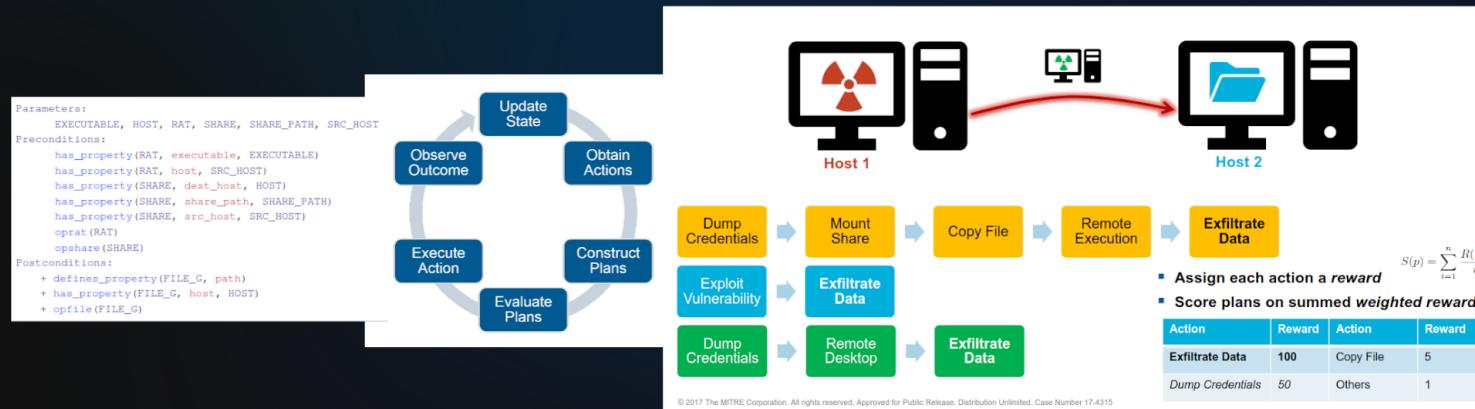
C2 Piloting / Heuristics

Caldera⁽¹⁾, an automated framework for adversary emulation (by Mitre), rely on a heuristic to select attacks & targets.



C2 Piloting / Heuristics

Caldera⁽¹⁾, an automated framework for adversary emulation (by Mitre), rely on a heuristic to select attacks & targets.



Based on pre & post conditions, really basic and not very "smart" system.

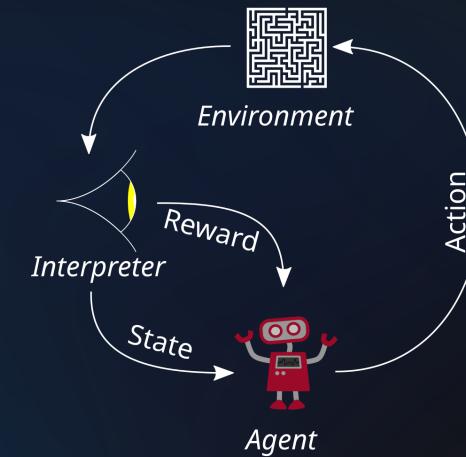
C2 Piloting / RL

C2 Piloting / RL

Another approach used by the literature is Reinforcement Learning (ML).

C2 Piloting / RL

Another approach used by the literature is Reinforcement Learning (ML).



C2 Piloting / RL

Another approach used by the literature is Reinforcement Learning (ML).

- S is the set of all states,
- A is the set of all actions,
- $R : S \times A \times S \rightarrow \mathbb{R}$ is the reward function,
- For a step k , $r_k = R(s_k, a_k, s_{k+1})$ is the associated reward function,

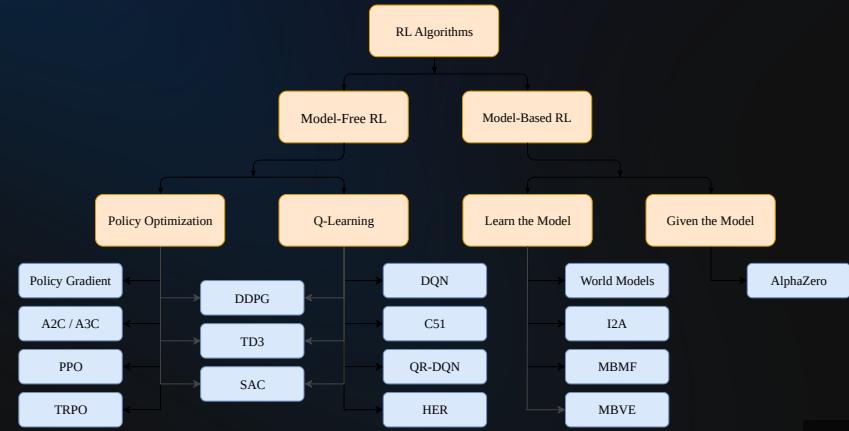
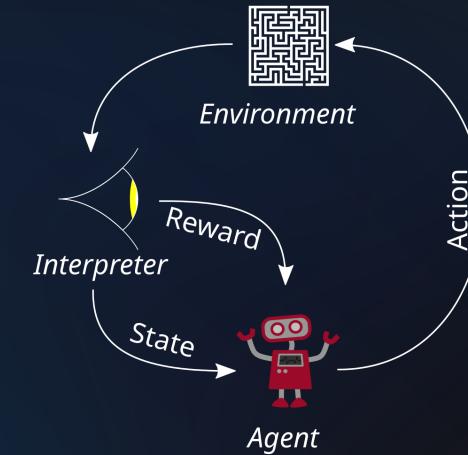


C2 Piloting / RL

Another approach used by the literature is Reinforcement Learning (ML).

- S is the set of all states,
- A is the set of all actions,
- $R : S \times A \times S \rightarrow \mathbb{R}$ is the reward function,
- For a step k , $r_k = R(s_k, a_k, s_{k+1})$ is the associated reward function,

Reinforcement learning is vast:



Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooo●oooooooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

C2 Piloting / RL for Piloting

C2 Piloting / RL for Piloting

RL fit well for C2 piloting:

- S contains all network states (like machines, OSes, open ports...),
- A contains all the modules in the attack library (**times possible parameters?!**),
- R is possible to construct, has demonstrated with Caldera's heuristic (**not that easy...**).

C2 Piloting / RL for Piloting

RL fit well for C2 piloting:

- S contains all network states (like machines, OSes, open ports...),
- A contains all the modules in the attack library (**times possible parameters?!**),
- R is possible to construct, has demonstrated with Caldera's heuristic (**not that easy...**).

Playing with R make it possible to complexify model behavior (e.g., valorize exploration vs exploitation, take defense into consideration...)

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooo●oooooooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

C2 Piloting / RL Environment

C2 Piloting / RL Environment

First issue, to use RL you need a training environment that'll simulate (mock) inputs/outputs ⁽²⁾⁽³⁾. This is called a simulation environment (or Gym).

Multiple environments are available, but are globally oriented toward Deep Learning, making them **very basic**:

C2 Piloting / RL Environment

First issue, to use RL you need a training environment that'll simulate (mock) inputs/outputs ⁽²⁾⁽³⁾. This is called a simulation environment (or Gym).

Multiple environments are available, but are globally oriented toward Deep Learning, making them **very basic**:

| Address | Compr. | Reach. | Disc. | Value | Access | linux | windows | proftpd | drupal | phpwiki | e_search | wp_ninja | mysql |
|---------|--------|--------|-------|-------|--------|-------|---------|---------|--------|---------|----------|----------|-------|
| (1, 0) | True | True | True | 0.0 | 1.0 | False | False | False | False | False | True | True | False |
| (3, 1) | False | True | True | 0.0 | 0.0 | False | False | False | False | False | False | False | False |
| (3, 0) | True | True | True | 100.0 | 2.0 | False | False | True | True | True | False | False | True |
| (4, 1) | False | True | True | 0.0 | 0.0 | False | False | False | False | False | False | False | False |
| (4, 0) | True | True | True | 0.0 | 1.0 | False | False | True | False | True | False | False | False |

(3)

C2 Piloting / RL Environment

First issue, to use RL you need a training environment that'll simulate (mock) inputs/outputs ⁽²⁾⁽³⁾. This is called a simulation environment (or Gym).

Multiple environments are available, but are globally oriented toward Deep Learning, making them **very basic**:

| Address | Compr. | Reach. | Disc. | Value | Access | linux | windows | proftpd | drupal | phpwiki | e_search | wp_ninja | mysql |
|---------|--------|--------|-------|-------|--------|-------|---------|---------|--------|---------|----------|----------|-------|
| (1, 0) | True | True | True | 0.0 | 1.0 | False | False | False | False | False | True | True | False |
| (3, 1) | False | True | True | 0.0 | 0.0 | False | False | False | False | False | False | False | False |
| (3, 0) | True | True | True | 100.0 | 2.0 | False | False | True | True | True | False | False | True |
| (4, 1) | False | True | True | 0.0 | 0.0 | False | False | False | False | False | False | False | False |
| (4, 0) | True | True | True | 0.0 | 1.0 | False | False | True | False | True | False | False | False |

(3)

Most, if not all, are **oversimplifying** the problem, and some doesn't even provide realistic evaluation (emulation, or running in a real IS). For example, none allow **attack parametrisation**. And most are using a probability of failure per action (artificial realism).

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooo●oooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

C2 Piloting / Some RL Environments

C2 Piloting / Some RL Environments

| Name | Release Date | Features |
|----------------|--------------|---|
| NASim | 2019 | Simulation only |
| CyberBattleSim | 2021 | Simulation only, Microsoft |
| CybORG (CAGE) | 2022 | Simulation & some level of emulation (~Metasploit C2) |
| NASimEmu | 2023 | Simulation & Emulation (Metasploit C2) |

C2 Piloting / Some RL Environments

| Name | Release Date | Features |
|----------------|--------------|---|
| NASim | 2019 | Simulation only |
| CyberBattleSim | 2021 | Simulation only, Microsoft |
| CybORG (CAGE) | 2022 | Simulation & some level of emulation (~Metasploit C2) |
| NASimEmu | 2023 | Simulation & Emulation (Metasploit C2) |

Writing such environment takes time since the **level of abstraction is very high**. There is also a **scaling issue**.

C2 Piloting / Some RL Environments

| Name | Release Date | Features |
|----------------|--------------|---|
| NASim | 2019 | Simulation only |
| CyberBattleSim | 2021 | Simulation only, Microsoft |
| CybORG (CAGE) | 2022 | Simulation & some level of emulation (~Metasploit C2) |
| NASimEmu | 2023 | Simulation & Emulation (Metasploit C2) |

Writing such environment takes time since the **level of abstraction is very high**. There is also a **scaling issue**.

Metasploit is the **ONLY** used C2 across the literature.

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooo●oooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

C2 Piloting / More RL Agents

C2 Piloting / More RL Agents

| Model | Technique | Complex Behavior | Parametrisation | Generalization | Year |
|------------------------------|-------------|---------------------|-----------------|----------------|------|
| POMDP ⁽⁴⁾ | POMDP | No | No | No | 2013 |
| LD-PenTesting ⁽⁵⁾ | POMDP | ~Defender Behaviour | No | No | 2020 |
| EPPTA ⁽⁶⁾ | POMDP + PPO | No | No | No | 2023 |
| AutoRed ⁽⁷⁾ | POMDP | No | No | GNN | 2024 |

C2 Piloting / More RL Agents

| Model | Technique | Complex Behavior | Parametrisation | Generalization | Year |
|------------------------------|-------------|---------------------|-----------------|----------------|------|
| POMDP ⁽⁴⁾ | POMDP | No | No | No | 2013 |
| LD-PenTesting ⁽⁵⁾ | POMDP | ~Defender Behaviour | No | No | 2020 |
| EPPTA ⁽⁶⁾ | POMDP + PPO | No | No | No | 2023 |
| AutoRed ⁽⁷⁾ | POMDP | No | No | GNN | 2024 |

Multiple issues with these approaches :

- Questionable generalization (except AutoRed),
- Imperfect reward function R /policy π (complexity),
- No attack parametrisation,
- Questionable performances,
- Low quality/high abstraction environments.

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooooooo●oooo

OAI for Evasion

oooooooooooooooooooo

Conclusion

oooo

C2 Piloting / A Complex Problem

C2 Piloting / A Complex Problem

Multiple questions:

- **One or multiple** models to select the target, the attack and the parameters?
- Does the target selection model **see the whole network**, or just select the target with the best attack?
- Since a network can have a (virtually) **infinite number of machines** (targets), how to construct the input/context for the target selection model?
- How to create an adequate **reward function/policy**?
- How to capture the **complexity** of the problem (like taking into account defense reaction)?
- C2 Piloting is a **planning problem** (with goals & subgoals)?
- On a big network, with a lot of possible actions, each with many parameters, how **not to collapse**?

Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○●○○

OAI for Evasion

○○○○○○○○○○○○○○○○

Conclusion

○○○○

C2 Piloting / New Approach

C2 Piloting / New Approach

Proposed approach

| Issue | Solution | Description |
|----------------------|------------------------------------|---|
| Generalization | GNN / Carrousel | Ingest the network context independently from its size. Learn to retain only interesting information. Context-based ingestion. |
| Partially Observable | POMDP + Dreamer (<i>JEPAs</i>) | Internal environment representation, latent space reasoning. |
| Complex Behavior | Complex <i>R</i> & Goals | Complexify the reward function. Add goals & subgoals to complexify policies (planning). |
| C2 Automation | New Specific C2 | Create a C2 with automation in mind. |
| RL Environment | New Emulation First RL Environment | Create a new RL environment centered around emulation (maximum realism). |

Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○○○●○

OAI for Evasion

○○○○○○○○○○○○○○○○

Conclusion

○○○○

C2 Piloting / Other Approaches

C2 Piloting / Other Approaches

Other approaches exist:

- Some people are trying to use **LLMs for C2 piloting** (me! ⁽⁸⁾), but also end-to-end OAI. Results are not great, OAI is not really **hallucination-compatible** (e.g., commands needs to be exact, IPs also, impact if the wrong target...).
- RL is not the only learning method: **Self-Supervised Learning** is trending.
- **Heuristics** are not dead.
- **Expert systems** (analogy to the defensive side) should also be explored.

C2 Piloting / Other Approaches

Other approaches exist:

- Some people are trying to use **LLMs for C2 piloting** (me! ⁽⁸⁾), but also end-to-end OAI. Results are not great, OAI is not really **hallucination-compatible** (e.g., commands needs to be exact, IPs also, impact if the wrong target...).
- RL is not the only learning method: **Self-Supervised Learning** is trending.
- **Heuristics** are not dead.
- **Expert systems** (analogy to the defensive side) should also be explored.

The dream is to have **end-to-end OAI** (target, attack, parameters choosing with attack generation).

Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooooooo●

OAI for Evasion

oooooooooooooooo

Conclusion

oooo

C2 Piloting / Conclusion

C2 Piloting / Conclusion

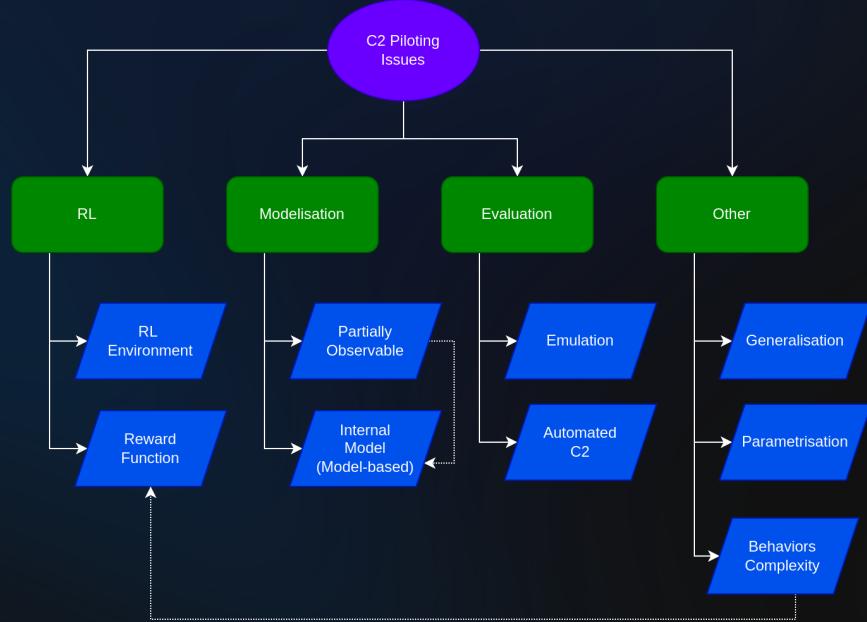
Some "Hot Takes":

- C2 Piloting is **hard**.
- This is not just a **model choice issue**.
- Past literature (minus some exception) are
focusing on the wrong problematic.
- Linked **use cases** are huge.
- Ask many **fundamental** questions.
- When it will work, the **impact might be important**
for defenses.
- *(LLMs are not the solution)*

C2 Piloting / Conclusion

Some "Hot Takes":

- C2 Piloting is hard.
- This is not just a **model choice issue**.
- Past literature (minus some exception) are **focusing on the wrong problematic**.
- Linked **use cases** are huge.
- Ask many **fundamental** questions.
- When it will work, the **impact might be important** for defenses.
- (*LLMs are not the solution*)



OAI for Evasion

Online & Offline evasion...

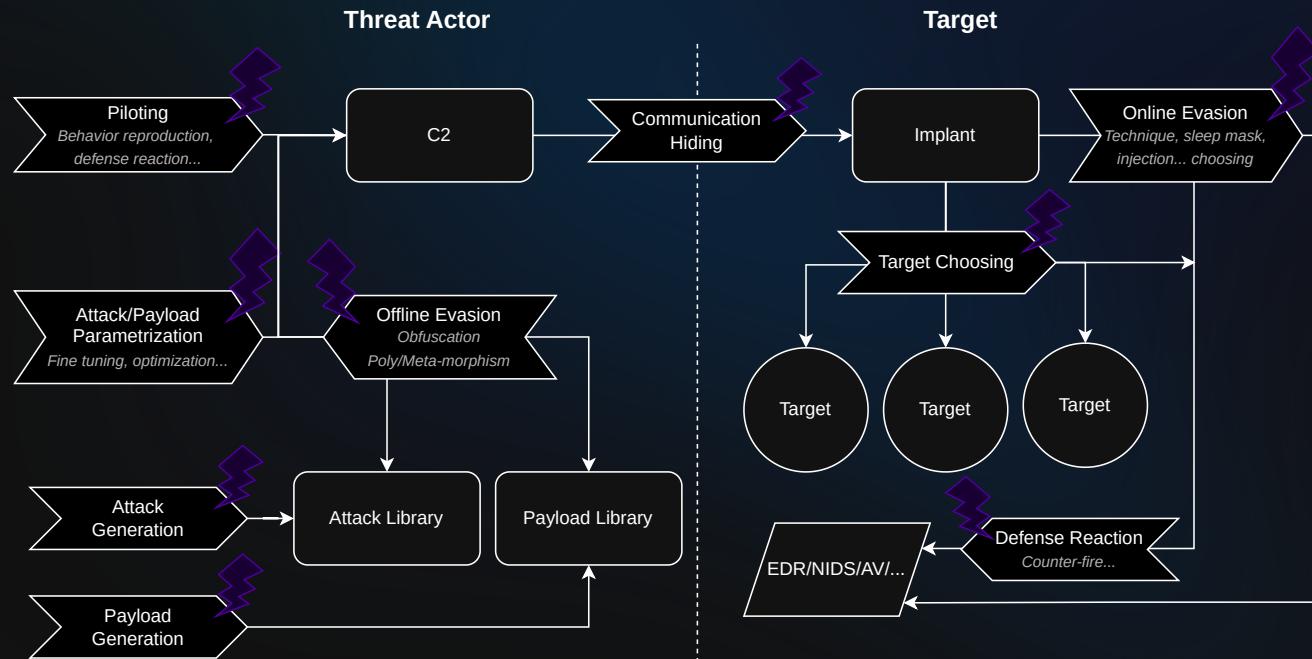
Evasion

Evasion

Consists of **optimizing** (e.g., by selecting parameters) an evasion method.

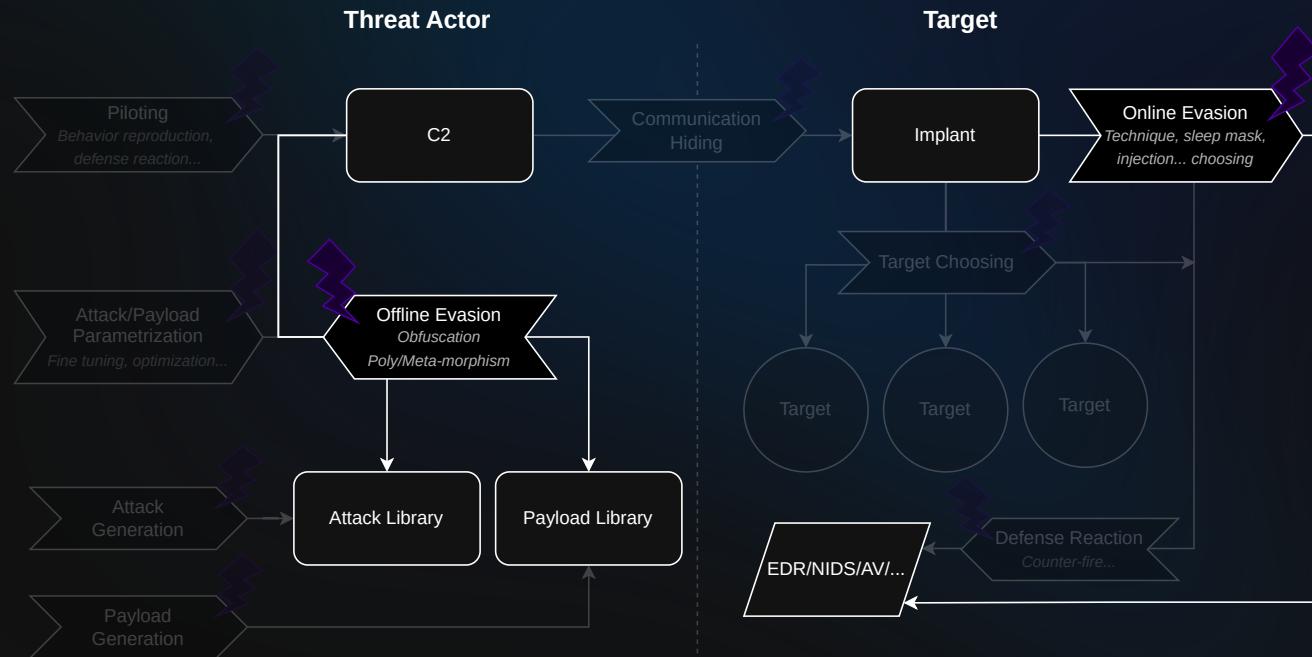
Evasion

Consists of optimizing (e.g., by selecting parameters) an evasion method.



Evasion

Consists of optimizing (e.g., by selecting parameters) an evasion method.



Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○○○○○○○

OAI for Evasion

○○●○○○○○○○○○○○○○○

Conclusion

○○○○

Evasion / The Idea

Evasion / The Idea

You want to **bypass defenses** (and here, specifically host based sensors, like EDRs).

Many evasion methods are available, but they generally have **parameters**.

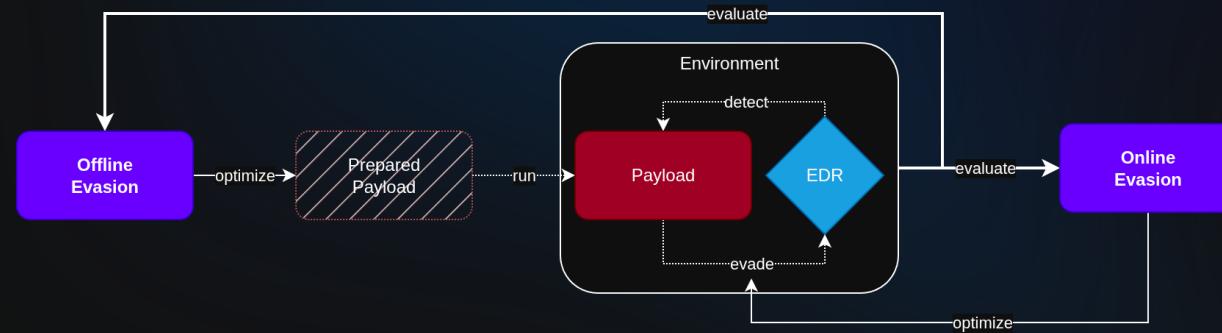
Each sensor is generally **more vulnerable** to specific combination of evasion techniques, associated with specific parameters.

Evasion / The Idea

You want to **bypass defenses** (and here, specifically host based sensors, like EDRs).

Many evasion methods are available, but they generally have **parameters**.

Each sensor is generally **more vulnerable** to specific combination of evasion techniques, associated with specific parameters.

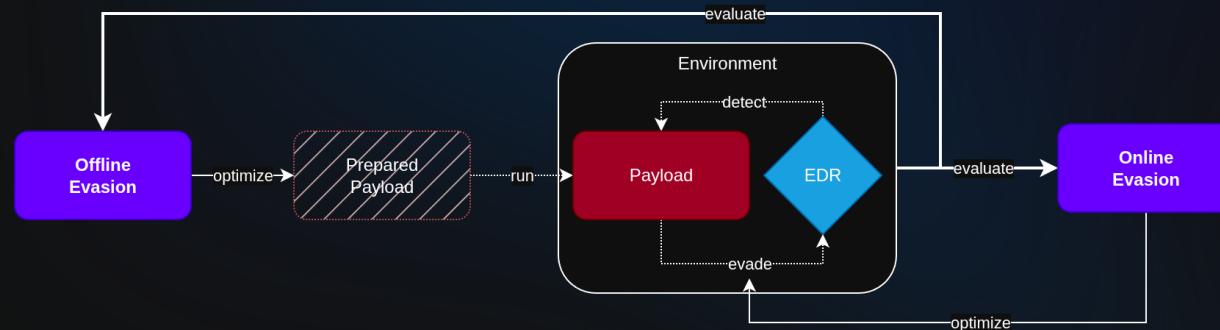


Evasion / The Idea

You want to **bypass defenses** (and here, specifically host based sensors, like EDRs).

Many evasion methods are available, but they generally have **parameters**.

Each sensor is generally **more vulnerable** to specific combination of evasion techniques, associated with specific parameters.



AI can be used to find these combinations and associated parameters, **optimizing evasion**.

Offline Evasion

Offline Evasion

Offline evasion includes any method that reduces detection & evade defenses **statically** (e.g., modifying a payload before sending it). We want to **obfuscate a payload** not to be detected by static analysis.

Offline Evasion

Offline evasion includes any method that reduces detection & evade defenses **statically** (e.g., modifying a payload before sending it). We want to **obfuscate a payload** not to be detected by static analysis.

This cover:

- Not having the same hash,
- Obfuscating code,
- Hidding strings,
- Packing,
- ...

Offline Evasion

Offline evasion includes any method that reduces detection & evade defenses **statically** (e.g., modifying a payload before sending it). We want to **obfuscate a payload** not to be detected by static analysis.

This cover:

- Not having the same hash,
- Obfuscating code,
- Hidding strings,
- Packing,
- ...

Some techniques:

- Polymorphism
- Metamorphism
- Obfuscation (excluding *morphism)
- ...

Offline Evasion

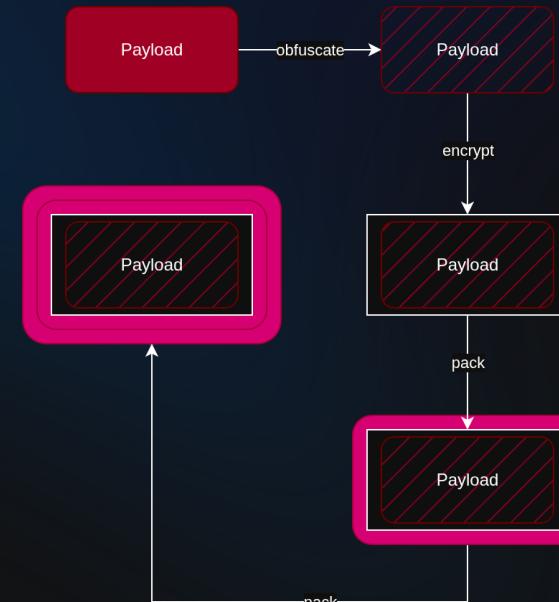
Offline evasion includes any method that reduces detection & evade defenses **statically** (e.g., modifying a payload before sending it). We want to obfuscate a payload not to be detected by static analysis.

This cover:

- Not having the same hash,
- Obfuscating code,
- Hidding strings,
- Packing,
- ...

Some techniques:

- Polymorphism
- Metamorphism
- Obfuscation (excluding *morphism)
- ...



Offline Evasion / *morphism

Offline Evasion / *morphism

Objective: Complexify part of a binary to make detection and/or reverse engineering harder. This maintains same functionalities, but make analysis more difficult.

Offline Evasion / *morphism

Objective: Complexify part of a binary to make detection and/or reverse engineering harder. This maintains same functionalities, but make analysis more difficult.

Automation Surface:

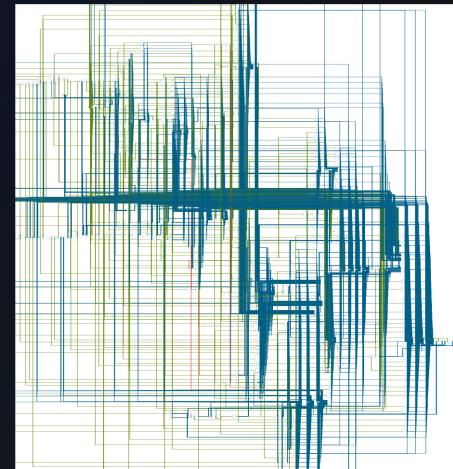
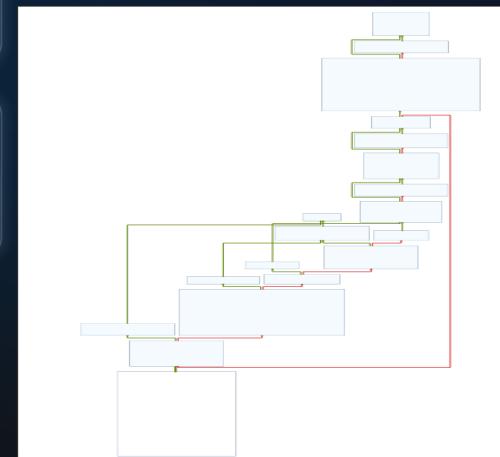
- Call graph obfuscation

Offline Evasion / *morphism

Objective: Complexify part of a binary to make detection and/or reverse engineering harder. This maintains same functionalities, but make analysis more difficult.

Automation Surface:

- Call graph obfuscation



Offline Evasion / *morphism

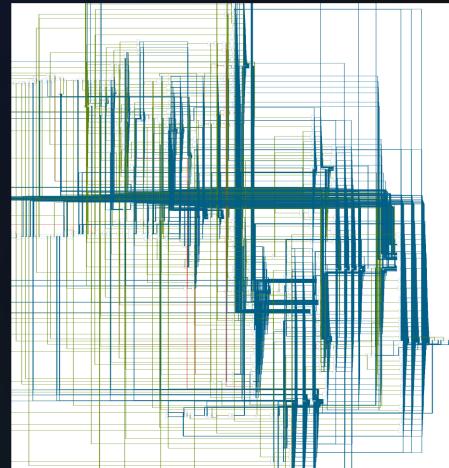
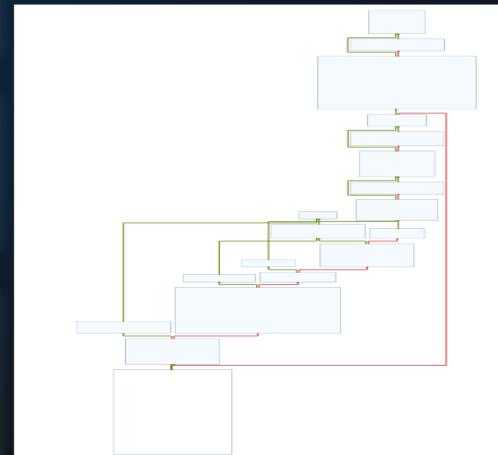
Objective: Complexify part of a binary to make detection and/or reverse engineering harder. This maintains same functionalities, but make analysis more difficult.

Automation Surface:

- Call graph obfuscation

Examples:

- 1 Maximize binary diff
- 2 Maximize number of subroutines
- 3 Maximize number of calls between each subroutines



Offline Evasion / *morphism

Objective: Complexify part of a binary to make detection and/or reverse engineering harder. This maintains same functionalities, but make analysis more difficult.

Automation Surface:

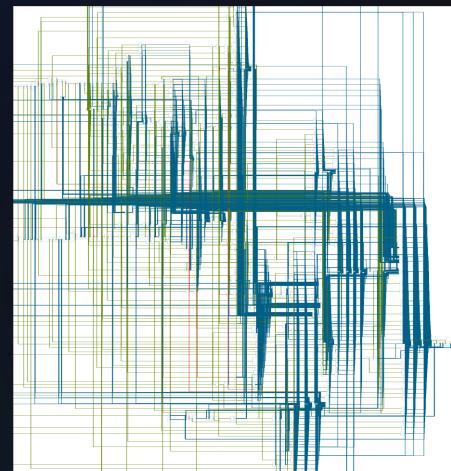
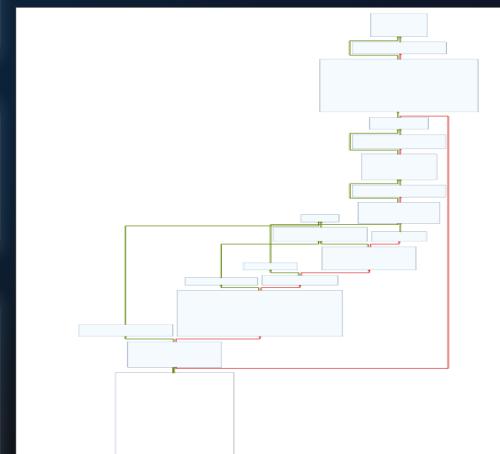
- Call graph obfuscation

Examples:

- 1 Maximize binary diff
- 2 Maximize number of subroutines
- 3 Maximize number of calls between each subroutines

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm): [1](#) [2](#) [3](#)
- Reinforcement Learning: [1](#) [2](#) [3](#)
- Heuristic: [2](#) [3](#)



Offline Evasion / Packing

Offline Evasion / Packing

Objective: Hide a payload by including it into another. At execution, the latter will **unpack** the payload and **pass execution** to it. Can include payload encryption & compression.

Offline Evasion / Packing

Objective: Hide a payload by including it into another. At execution, the latter will **unpack** the payload and **pass execution** to it. Can include payload encryption & compression.

Automation Surface:

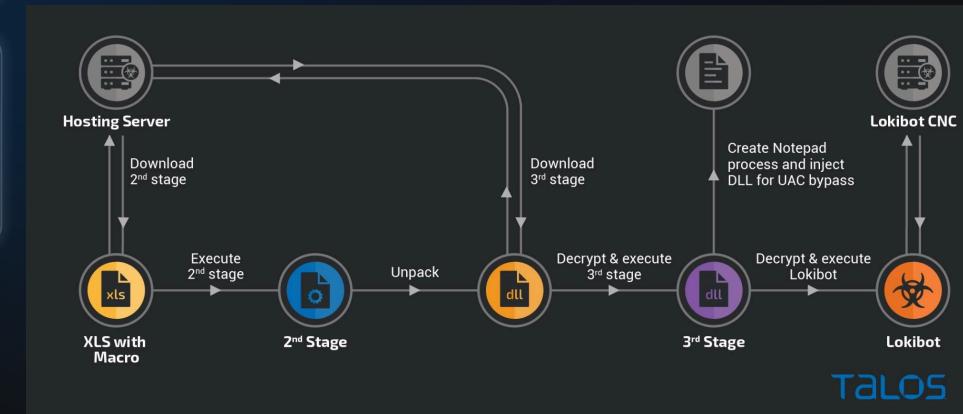
- Packer nesting
- Packer optimization

Offline Evasion / Packing

Objective: Hide a payload by including it into another. At execution, the latter will **unpack** the payload and **pass execution** to it. Can include payload encryption & compression.

Automation Surface:

- o Packer nesting
- o Packer optimization



Offline Evasion / Packing

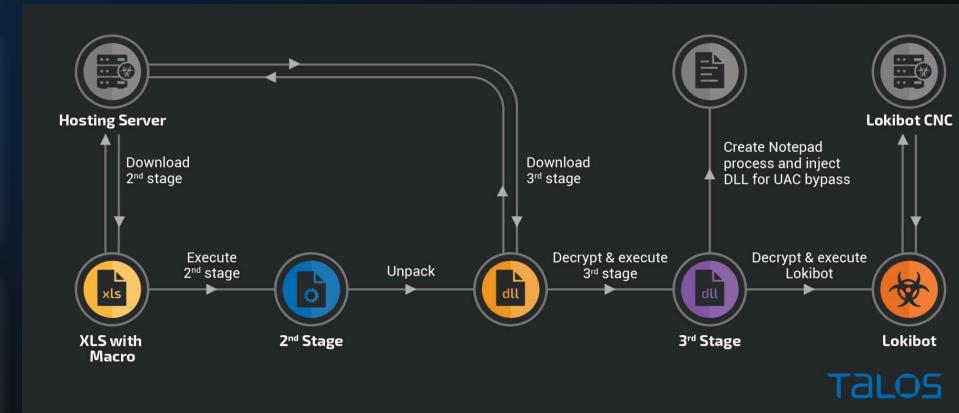
Objective: Hide a payload by including it into another. At execution, the latter will **unpack** the payload and **pass execution** to it. Can include payload encryption & compression.

Automation Surface:

- Packer nesting
- Packer optimization

Examples:

- 1 Minimize packer entropy
- 2 Maximize waves diversity



TALOS

Offline Evasion / Packing

Objective: Hide a payload by including it into another. At execution, the latter will **unpack** the payload and **pass execution** to it. Can include payload encryption & compression.

Automation Surface:

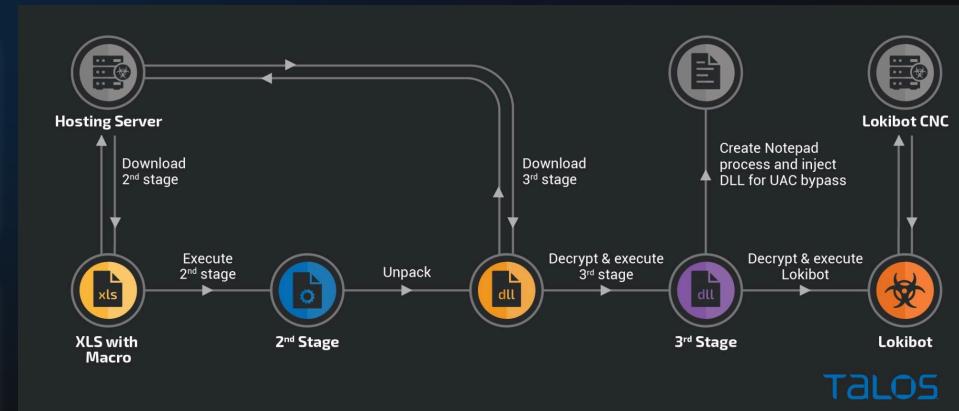
- Packer nesting
- Packer optimization

Examples:

- 1 Minimize packer entropy
- 2 Maximize waves diversity

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm) 1 2
- Reinforcement Learning 1 2
- Heuristic: 1



Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○○○○○○○

OAI for Evasion

○○○○○●○○○○○○○○

Conclusion

○○○○

Online Evasion

Online Evasion

Definition

Online evasion includes any method that **reduce detection & evade** defenses **dynamically** (at runtime). We want to **hide traces** not to be detected by dynamic analysis.

Online Evasion

Definition

Online evasion includes any method that **reduce detection & evade** defenses **dynamically** (at runtime). We want to **hide traces** not to be detected by dynamic analysis.

This cover:

- Hide API calls,
- Hide in another process,
- Mask/Encrypt part of the payload in memory,
- Reproduce legitimate behaviors,
- ...

Online Evasion

Definition

Online evasion includes any method that **reduce detection & evade** defenses **dynamically** (at runtime). We want to **hide traces** not to be detected by dynamic analysis.

This cover:

- Hide API calls,
- Hide in another process,
- Mask/Encrypt part of the payload in memory,
- Reproduce legitimate behaviors,
- ...

Some techniques:

- Sleep Obfuscation
- Memory Masking
- Process Injection
- Call Stack Spoofing
- ...

Online Evasion

Definition

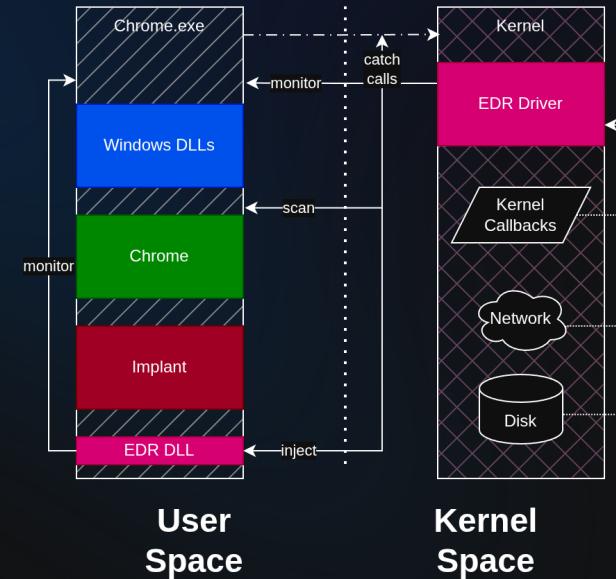
Online evasion includes any method that **reduce detection & evade defenses dynamically** (at runtime). We want to **hide traces** not to be detected by dynamic analysis.

This cover:

- Hide API calls,
- Hide in another process,
- Mask/Encrypt part of the payload in memory,
- Reproduce legitimate behaviors,
- ...

Some techniques:

- Sleep Obfuscation
- Memory Masking
- Process Injection
- Call Stack Spoofing
- ...



Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooooooooooooooo

OAI for Evasion

oooooooo●oooooooo

Conclusion

oooo

Online Evasion / Process Injection

Online Evasion / Process Injection

Objective: Execute malicious code in the context of a **target process**.

This can turn any legitimate process into a malicious one.

Online Evasion / Process Injection

Objective: Execute malicious code in the context of a **target process**.

This can turn any legitimate process into a malicious one.

Automation Surface:

- Find API combination
- Choose target process

Online Evasion / Process Injection

Objective: Execute malicious code in the context of a **target process**.

This can turn any legitimate process into a malicious one.

Automation Surface:

- Find API combination
- Choose target process

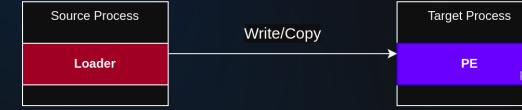
Step 1



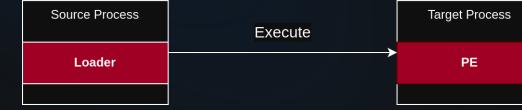
Step 2



Step 3



Step 4



Online Evasion / Process Injection

Objective: Execute malicious code in the context of a **target process**.

This can turn any legitimate process into a malicious one.

Automation Surface:

- Find API combination
- Choose target process

Examples:

- 1 Maximize behavior similarity of target process
- 2 Select best injection method (against specific defense)

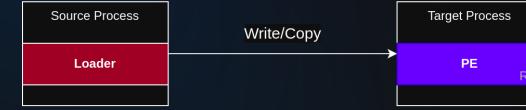
Step 1



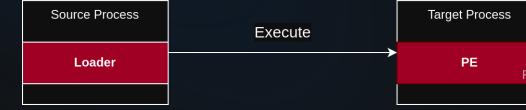
Step 2



Step 3



Step 4



Online Evasion / Process Injection

Objective: Execute malicious code in the context of a **target process**.

This can turn any legitimate process into a malicious one.

Automation Surface:

- Find API combination
- Choose target process

Examples:

- 1 Maximize behavior similarity of target process
- 2 Select best injection method (against specific defense)

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm): 1 2
- Reinforcement Learning: 1 2
- Heuristic: 1

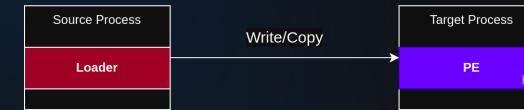
Step 1



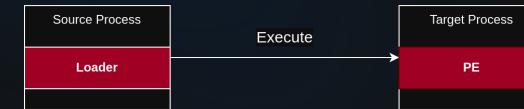
Step 2



Step 3



Step 4



Online Evasion / Memory Masking

Online Evasion / Memory Masking

Objective: Try to **hide** part, or the entirety, of a payload **in memory**.

This is useful against **memory scanners**, like YARA. **Sleep**

obfuscation is a variant, where a payload will sleep & be encrypted
in memory during that time.

Online Evasion / Memory Masking

Objective: Try to **hide** part, or the entirety, of a payload **in memory**.

This is useful against **memory scanners**, like YARA. **Sleep**

obfuscation is a variant, where a payload will sleep & be encrypted in memory during that time.

Automation Surface:

- Technique parameters
- API used

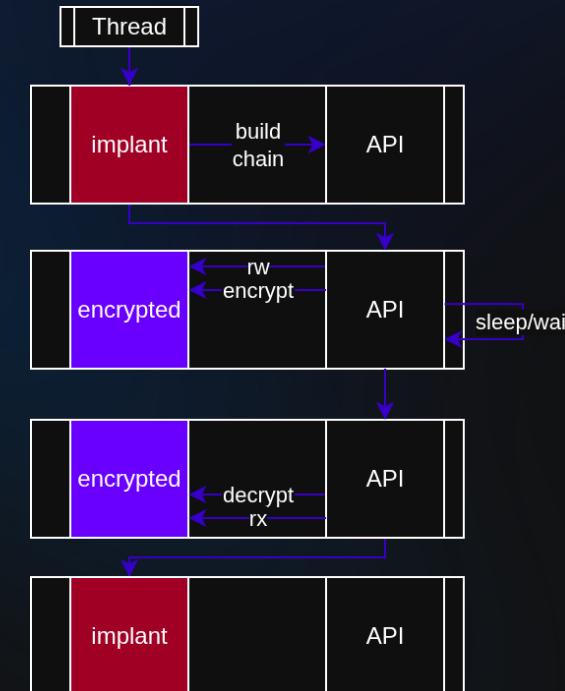
Online Evasion / Memory Masking

Objective: Try to **hide** part, or the entirety, of a payload **in memory**.

This is useful against **memory scanners**, like YARA. **Sleep obfuscation** is a variant, where a payload will sleep & be encrypted in memory during that time.

Automation Surface:

- Technique parameters
- API used



Online Evasion / Memory Masking

Objective: Try to **hide** part, or the entirety, of a payload **in memory**.

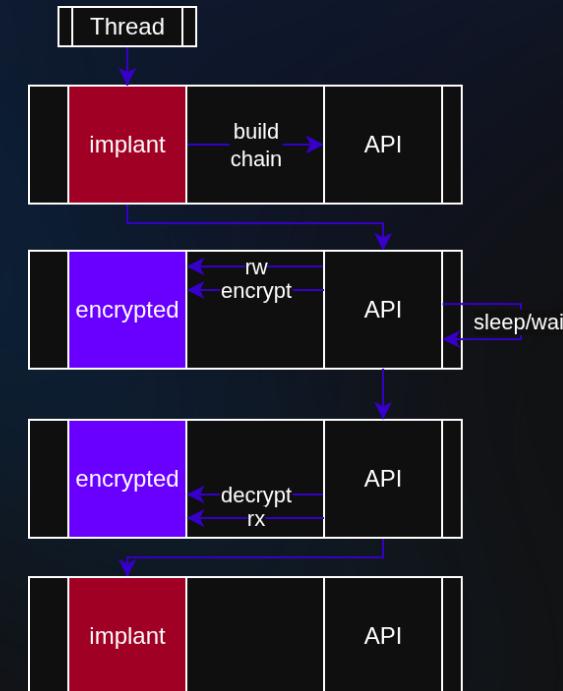
This is useful against **memory scanners**, like YARA. **Sleep obfuscation** is a variant, where a payload will sleep & be encrypted in memory during that time.

Automation Surface:

- Technique parameters
- API used

Examples:

- 1 Find "optimal" sleep/jitter value
- 2 Find "optimal" delay before remasking



Online Evasion / Memory Masking

Objective: Try to **hide** part, or the entirety, of a payload **in memory**.

This is useful against **memory scanners**, like YARA. **Sleep obfuscation** is a variant, where a payload will sleep & be encrypted in memory during that time.

Automation Surface:

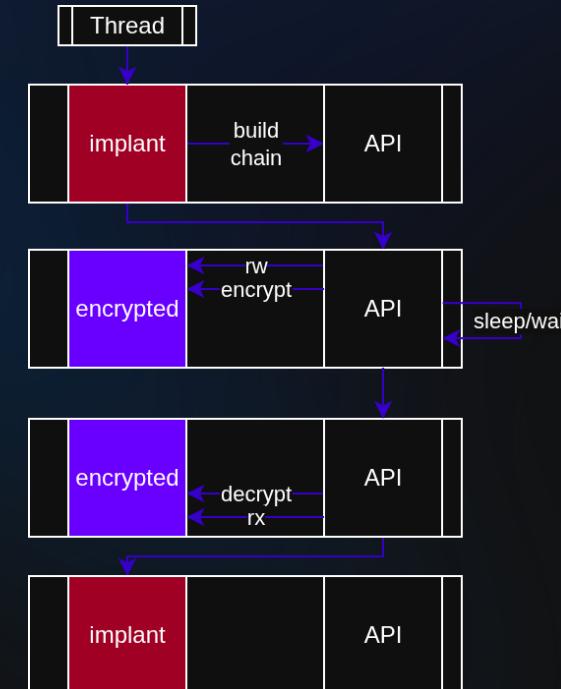
- Technique parameters
- API used

Examples:

- 1 Find "optimal" sleep/jitter value
- 2 Find "optimal" delay before remasking

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm): 1 2
- Heuristic: 1 2



Introduction

oooo

OAI

ooo

OAI for C2 Piloting

oooooooooooooooooooo

OAI for Evasion

oooooooo●ooooo

Conclusion

oooo

Online Evasion / Call Stack

Online Evasion / Call Stack

Objective: Spoof or construct a fake call stack, to hide original callers and evade detection.

Online Evasion / Call Stack

Objective: Spoof or construct a fake call stack, to hide original callers and evade detection.

Automation Surface:

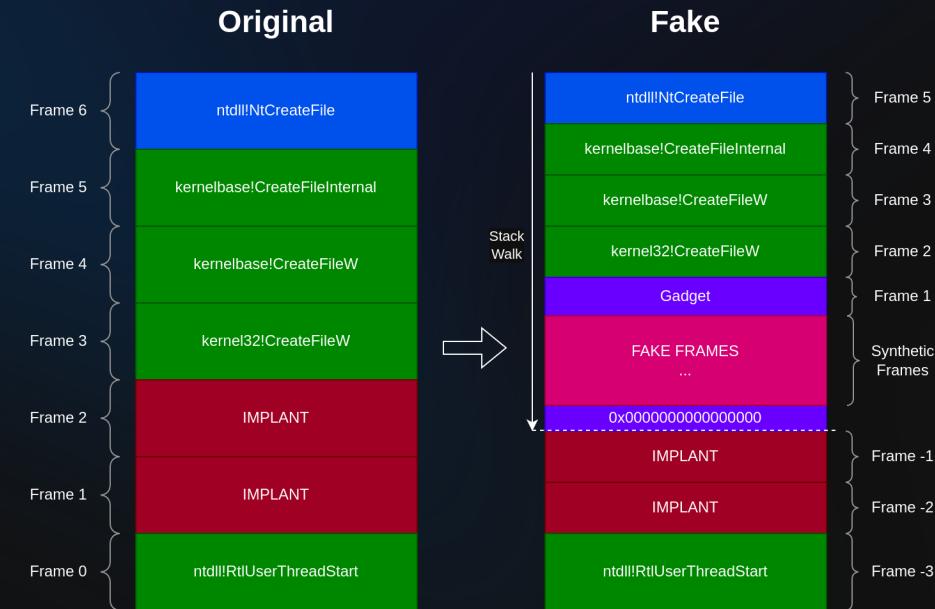
- Fake API choice

Online Evasion / Call Stack

Objective: Spoof or construct a fake call stack, to hide original callers and evade detection.

Automation Surface:

- o Fake API choice



Online Evasion / Call Stack

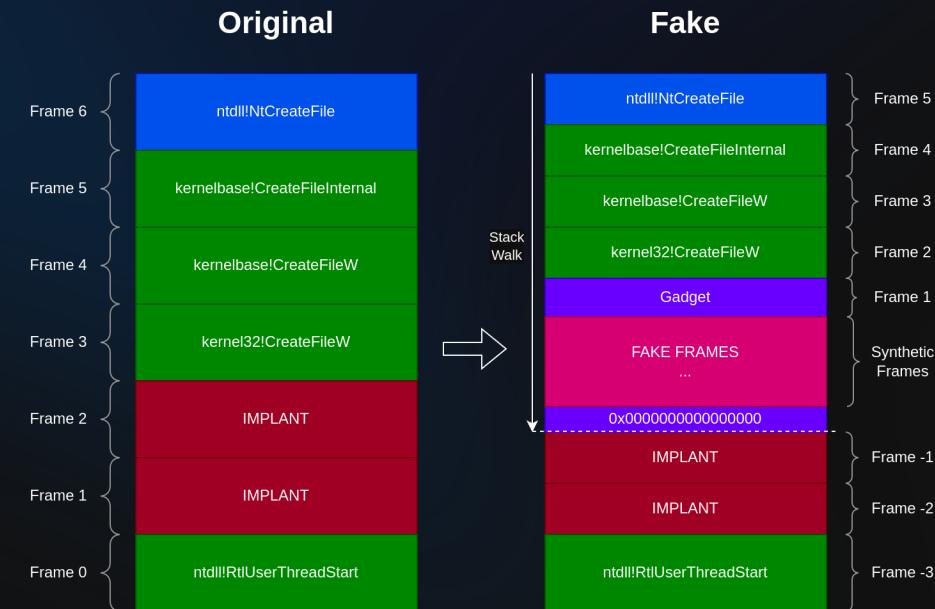
Objective: Spoof or construct a fake call stack, to hide original callers and evade detection.

Automation Surface:

- Fake API choice

Examples:

- 1 Maximize similarity to injected process call stack
- 2 Find suspicious call stacks



Online Evasion / Call Stack

Objective: Spoof or construct a fake call stack, to hide original callers and evade detection.

Automation Surface:

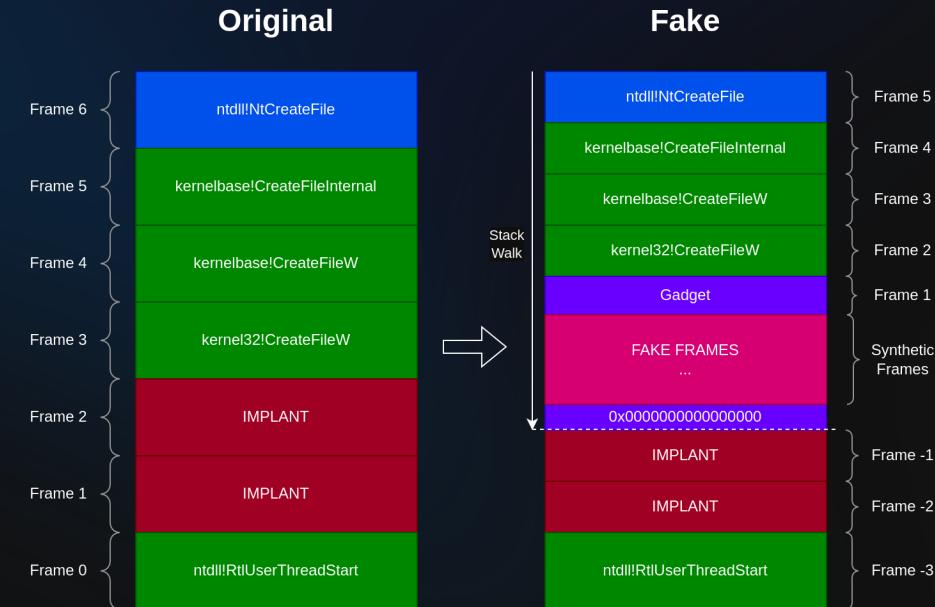
- Fake API choice

Examples:

- 1 Maximize similarity to injected process call stack
- 2 Find suspicious call stacks

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm): 1 2
- Reinforcement Learning: 1 2
- Heuristic: 1



Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○○○○○○○

OAI for Evasion

○○○○○○○○●○○○○

Conclusion

○○○○

Meta-optimization

Meta-optimization

Objective: Each evasion technique has some variation available.

Considering all variations of all techniques, this creates a lot of combination possibilities.

Meta-optimization

Objective: Each evasion technique has some variation available.

Considering all variations of all techniques, this creates a lot of combination possibilities.

Automation Surface:

- Techniques combinations
- Techniques variations

Meta-optimization

Objective: Each evasion technique has some variation available.

Considering all variations of all techniques, this creates a lot of combination possibilities.

Automation Surface:

- Techniques combinations
- Techniques variations

Examples:

- 1 Minimize number of techniques
- 2 Find all combinations leading to bypass

Meta-optimization

Objective: Each evasion technique has some variation available.

Considering all variations of all techniques, this creates a lot of combination possibilities.

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm): [1](#) [2](#)
- Reinforcement Learning: [1](#) [2](#)
- Surrogate Model: [2](#)

Automation Surface:

- Techniques combinations
- Techniques variations

Examples:

- [1](#) Minimize number of techniques
- [2](#) Find all combinations leading to bypass

Meta-optimization

Objective: Each evasion technique has some variation available. Considering all variations of all techniques, this creates a lot of combination possibilities.

Automation Surface:

- Techniques combinations
- Techniques variations

Examples:

- 1 Minimize number of techniques
- 2 Find all combinations leading to bypass

Automation Possibilities:

- Meta-heuristic (e.g., Genetic Algorithm): 1 2
- Reinforcement Learning: 1 2
- Surrogate Model: 2

| Product | Call Stack | Process Injection | Metamorphism | Memory Masking |
|--------------|------------|-------------------|--------------|----------------|
| EDR1 | Variant 2 | "Explorer" | Variant 1 | X |
| EDR2 | Variant 2 | X | Variant 1 | X |
| EDR2 + Win11 | Variant 3 | "Google Chrome" | X | Variant 2 |

Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○○○○○○○

OAI for Evasion

○○○○○○○○○○●○○○

Conclusion

○○○○

Surrogate Model

Surrogate Model

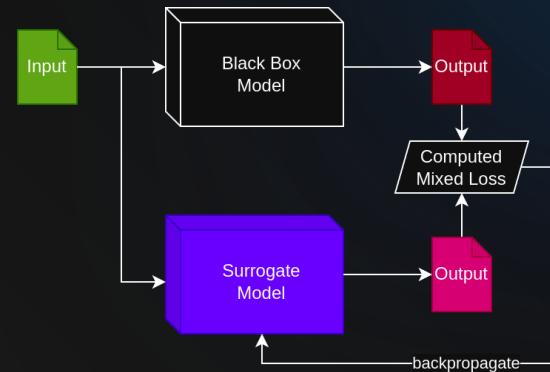
Definition

A surrogate model approximates a black box model by learning pairs of input/output.

Surrogate Model

Definition

A surrogate model approximates a black box model by learning pairs of input/output.



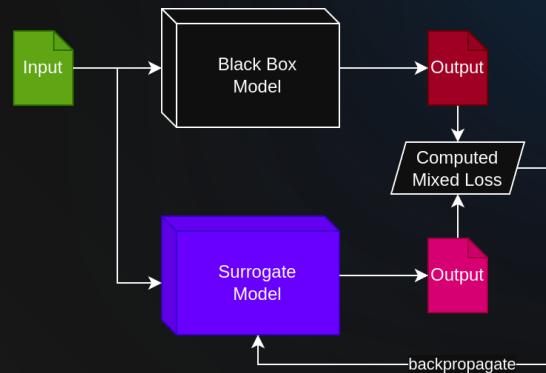
Surrogate Model

Definition

A surrogate model approximates a black box model by learning pairs of input/output.

Objective

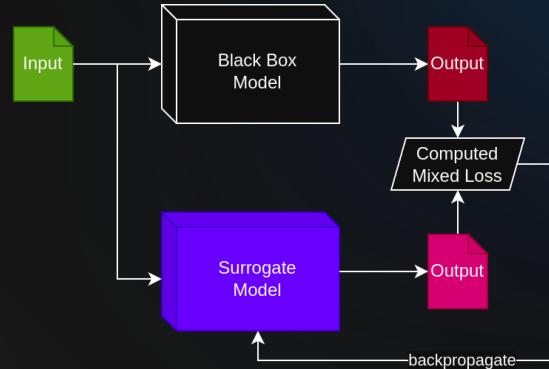
We can then use the surrogate model to build adversarial inputs to force misclassification by the black box model. For that, we can use the surrogate model internals (like embedding or gradients).



Surrogate Model

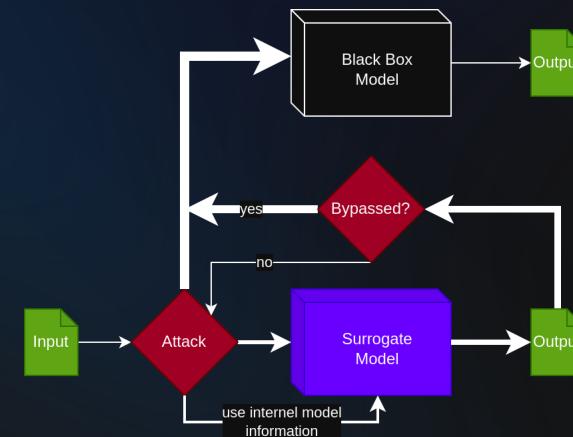
Definition

A surrogate model approximates a black box model by learning pairs of input/output.



Objective

We can then use the surrogate model to build adversarial inputs to force misclassification by the black box model. For that, we can use the surrogate model internals (like embedding or gradients).



Surrogate Model: Evasion

Surrogate Model: Evasion

Usage for evasion

The black box model can be the classification algorithm of an EDR, the input a binary (offline evasion) or dynamic traces generated by it (online evasion) and the output the benign/malicious classification.

Surrogate Model: Evasion

Usage for evasion

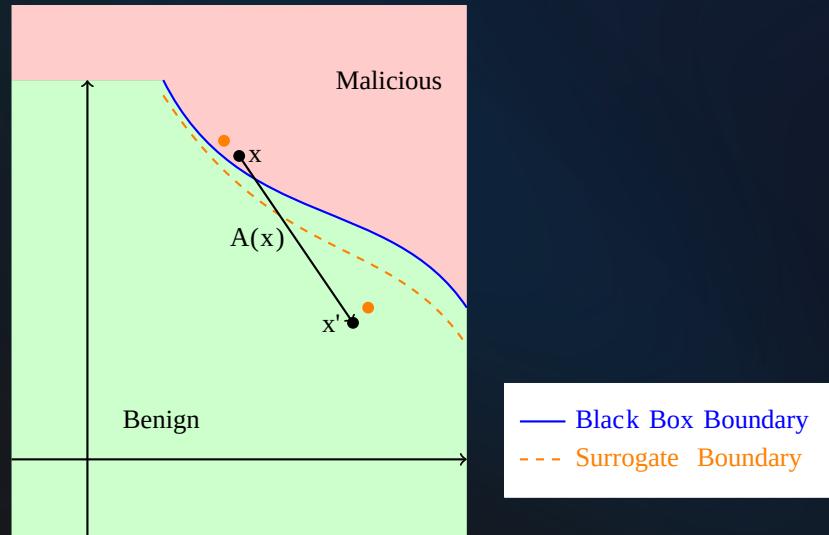
The black box model can be the classification algorithm of an EDR, the input a binary (offline evasion) or dynamic traces generated by it (online evasion) and the output the benign/malicious classification.

Possible attacks

- Add obfuscation, packing...
- Use process injection, call stack spoofing...
- Mask/Obfuscate memory...

All of these techniques are primitives that we can leverage to build adversarial binaries, using our surrogate model.

Surrogate Model: Visualization



Where $A(x) = \epsilon_\Omega(x)$.

Introduction

○○○○

OAI

○○○

OAI for C2 Piloting

○○○○○○○○○○○○○○○○

OAI for Evasion

○○○○○○○○○○○○○●

Conclusion

○○○○

Conclusion

Conclusion

Some "Hot Takes":

- Automation for evasion is **easier** than C2 piloting.
- Not much literature on the subject.
- Might not be explored by anyone?
- Multimodal > Unimodal
- Linked **use cases** are important.
- When it will work, the **impact might be important** for defenses.
- *(LLMs are not the solution)*

Conclusion

OAI is larger!

OAI is larger!

Not discussed, but:

- OAI for **attack generation**.
- OAI for **phishing**.
- OAI for **defense reaction**.
- OAI for **payload generation**.
- OAI for **communication hiding**.
- ...

OAI is larger!

Not discussed, but:

- OAI for attack generation.
 - OAI for phishing.
 - OAI for defense reaction.
 - OAI for payload generation.
 - OAI for communication hiding.
 - ...

The subject is vast!

OAI is larger!

Not discussed, but:

- OAI for **attack generation**.
- OAI for **phishing**.
- OAI for **defense reaction**.
- OAI for **payload generation**.
- OAI for **communication hiding**.
- ...

The subject is **vast!**

But **not very explored** (in the literature).

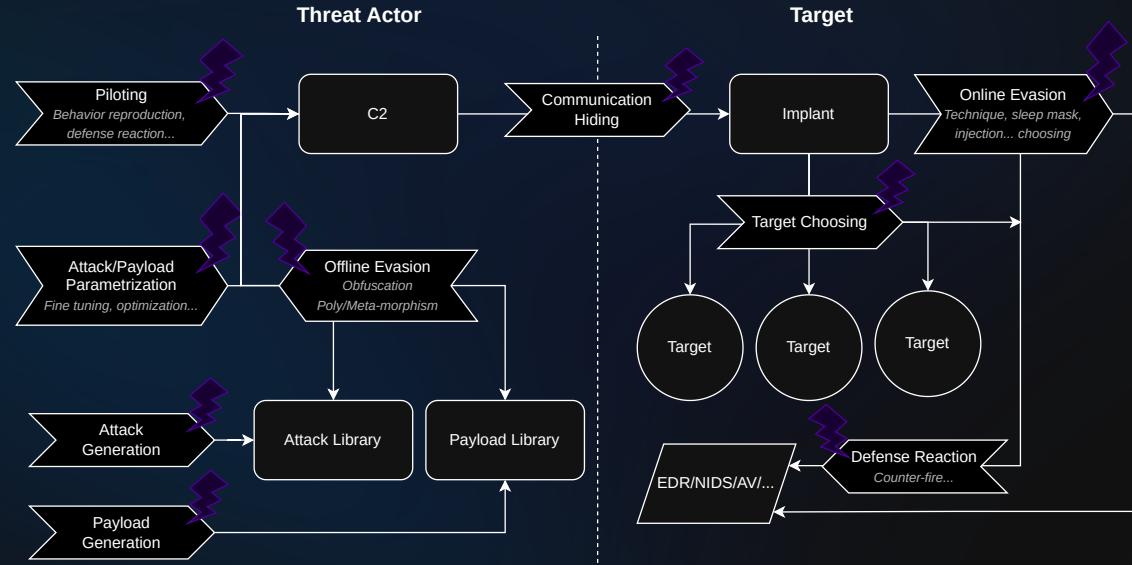
OAI is larger!

Not discussed, but:

- OAI for **attack generation**.
- OAI for **phishing**.
- OAI for **defense reaction**.
- OAI for **payload generation**.
- OAI for **communication hiding**.
- ...

The subject is vast!

But not very explored (in the literature).



Thanks!

A complete version of this conference is available here ↗ <https://dorianb.net/talks/OAI2026/Full>

For more ↗ <https://dorianb.net/>

Bibliography

- 1. Miller D, Alford R, Applebaum A, Foster H, Little C, Strom BE. Automated Adversary Emulation: A Case for Planning and Acting with Unknowns. In 2018. Available from: <https://api.semanticscholar.org/CorpusID:145814722> (13)
- 2. Schwartz J, Kurniawati H. Autonomous Penetration Testing using Reinforcement Learning. ArXiv [Internet]. 2019;abs/1905.05965. Available from: <https://api.semanticscholar.org/CorpusID:155090057> (16)
- 3. Janisch J, Pevn'y T, Lis'y V. NASimEmu: Network Attack Simulator & Emulator for Training Agents Generalizing to Novel Scenarios. ArXiv [Internet]. 2023;abs/2305.17246. Available from: <https://api.semanticscholar.org/CorpusID:258959386> (16)

page 1 / 5

Bibliography

- 4. Sarraute C, Buffet O, Hoffmann J. POMDPs Make Better Hackers: Accounting for Uncertainty in Penetration Testing [Internet]. 2013. Available from: <https://arxiv.org/abs/1307.8182> (18)
- 5. Schwartz J, Kurniawati H, El-Mahassni E. POMDP + Information-Decay: Incorporating Defender's Behaviour in Autonomous Penetration Testing. Proceedings of the International Conference on Automated Planning and Scheduling [Internet]. 2020;30(1):235–43. Available from: <https://ojs.aaai.org/index.php/ICAPS/article/view/6666> (18)
- 6. Li Z, Zhang Q, Yang G. EPPTA: Efficient partially observable reinforcement learning agent for penetration testing applications. Engineering Reports. 2023; (18)

Bibliography

- 7. Hasegawa K, Hidano S, Fukushima K. AutoRed: Automating Red Team Assessment via Strategic Thinking Using Reinforcement Learning. In: Proceedings of the Fourteenth ACM Conference on Data and Application Security and Privacy [Internet]. Porto, Portugal: Association for Computing Machinery; 2024. p. 325–36. (CODASPY '24). Available from: <https://doi.org/10.1145/3626232.3653252> (18)
- 8. Genevey-Metat C, Bachelot D, Gourmelen T, Quemat A, Satre PM, Scotto L, et al. Red Team LLM: towards an adaptive and robust automation solution. In: Conference on Artificial Intelligence for Defense [Internet]. Rennes, France: DGA Maîtrise de l'Information; 2023. Available from: <https://hal.science/hal-04328468> (21)
- 9. Zhou S, Liu J, Hou D, Zhong X, Zhang Y. Autonomous Penetration Testing Based on Improved Deep Q-Network. Applied Sciences [Internet]. 2021;11(19). Available from: <https://www.mdpi.com/2076-3417/11/19/8823>

Bibliography

- 10. Yang Y, Liu X. Behaviour-Diverse Automatic Penetration Testing: A Curiosity-Driven Multi-Objective Deep Reinforcement Learning Approach [Internet]. 2022. Available from: <https://arxiv.org/abs/2202.10630>
- 11. Blumenthal O, Shani G. Domain Independent Heuristics for Online Stochastic Contingent Planning. 2023.
- 12. Tran K, Akella A, Standen M, Kim J, Bowman D, Richer T, et al. Deep hierarchical reinforcement agents for automated penetration testing [Internet]. 2021. Available from: <https://arxiv.org/abs/2109.06449>

page 4 / 5

Bibliography

- 13. Chen Z, Kang F, Xiong X, Shu H. A Survey on Penetration Path Planning in Automated Penetration Testing. Applied Sciences [Internet]. 2024;14(18). Available from: <https://www.mdpi.com/2076-3417/14/18/8355>

page 5 / 5